

//Youtube Video:

--My Youtube Channel: [youtube.com/c/joycekayavila](https://www.youtube.com/c/joycekayavila)

--Link to the specific video associated with this tutorial: https://youtu.be/zb89rvl4_Xs

//PREWORK - ENVIRONMENT SETUP

//Let's set the context for this worksheet

//Note: for simplicity, the ACCOUNTADMIN role will be used for this tutorial.

--but this would not likely be the case in a real production environment.

USE ROLE ACCOUNTADMIN;

USE WAREHOUSE COMPUTE_WH;

//At the time of this tutorial, there were 15 million records in this table

SELECT COUNT(1) FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF10.ORDERS;

//At the time of this tutorial, there were 1.5 million records in this table

SELECT COUNT(1) FROM

SNOWFLAKE_SAMPLE_DATA.TPCH_SF10.CUSTOMER;

//Create two virtual warehouses, one for generic use

--and the other for use in the raw data layer for data pipelines.

//Note: To make it easy to go back to any section of the tutorial,

--we'll use the 'CREATE OR REPLACE' command to create new objects.

--The 'CREATE OR REPLACE' command is not recommended for common use in the production environment.

//Note: We are creating two new virtual warehouses but you can create as many virtual warehouses,

--in varying sizes and configurations, that you need.

//Tip: You may want to consider creating separate virtual warehouses for different layers

--in the Data Vault architecture such as DV_BDV_WH and DV_INFO_WH

CREATE OR REPLACE WAREHOUSE DV_GENERIC_WH

```
    WITH WAREHOUSE_SIZE = 'XSMALL' AUTO_SUSPEND = 300
    AUTO_RESUME = TRUE INITIALLY_SUSPENDED = TRUE;
CREATE OR REPLACE WAREHOUSE DV_RDV_WH
    WITH WAREHOUSE_SIZE = 'XSMALL' AUTO_SUSPEND = 300
    AUTO_RESUME = TRUE INITIALLY_SUSPENDED = TRUE;
```

```
//Use the new Generic virtual warehouse now for the tutorial
USE WAREHOUSE DV_GENERIC_WH;
```

```
//Create new Snowflake database and schemas to be used in the tutorial
CREATE OR REPLACE DATABASE DV_TUTORIAL;
CREATE OR REPLACE SCHEMA L00_STG COMMENT = 'Schema for Staging
Area objects';
CREATE OR REPLACE SCHEMA L10_RDV COMMENT = 'Schema for Raw Data
Vault objects';
CREATE OR REPLACE SCHEMA L20_BDV COMMENT = 'Schema for Business
Data Vault objects';
CREATE OR REPLACE SCHEMA L30_INFO COMMENT = 'Schema for Information
Delivery objects';
```

```
-----
-----
-----
```

```
//STAGING AREA SETUP - NATION & REGION
```

```
//Set Context. Use the schema for Staging area objects - L00.STG
USE SCHEMA L00_STG;
```

```
//Create two new staging tables for static reference data
--LDTS = Load Data Timestamp
--RSCR = Reference Source (Static Reference Data)
```

```
//Create Nation Stage Table
CREATE OR REPLACE TABLE STG_NATION
AS
SELECT SRC.*,
    CURRENT_TIMESTAMP() LDTS,
    'STATIC REFERENCE DATA' RSCR
FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF10.NATION SRC;
```

```
//Validate
```

```
SELECT * FROM STG_NATION;
```

```
//Create Region Table
```

```
CREATE OR REPLACE TABLE STG_REGION  
AS
```

```
SELECT SRC.*,  
       CURRENT_TIMESTAMP() LDTS,  
       'STATIC REFERENCE DATA' RSCR  
FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF10.REGION SRC;
```

```
//Validate
```

```
SELECT * FROM STG_REGION;
```

```
//STAGING AREA SETUP - CUSTOMER & ORDER
```

```
//Create two new tables to be used by Snowpipe to  
--drip-feed the data as it lands in the stage
```

```
//Note that the full payload of JSON data will be loaded into the raw_json  
--column. We'll use the special VARIANT data type for this use case
```

```
//Note that we'll also add some columns for metadata like  
--load data timestamp (ldts) and file row number
```

```
//Create the Customer Stage Table
```

```
CREATE OR REPLACE TABLE STG_CUSTOMER  
(RAW_JSON    VARIANT,  
 FILENAME    STRING NOT NULL,  
 FILE_ROW_SEQ NUMBER NOT NULL,  
 LDTS        STRING NOT NULL,  
 RSCR        STRING NOT NULL);
```

```
//No records yet entered so query produces no results
```

```
SELECT * FROM STG_CUSTOMER;
```

```
//Create the ORDER Stage Table
```

```
CREATE OR REPLACE TABLE STG_ORDER  
(O_ORDERKEY    NUMBER,  
 O_CUSTKEY     NUMBER,  
 O_ORDERSTATUS STRING,  
 O_TOTALPRICE  NUMBER,  
 O_ORDERDATE   DATE,  
 O_ORDERPRIORITY STRING,  
 O_CLERK       STRING,
```

```
O_SHIPPRIORITY    NUMBER,  
O_COMMENT         STRING,  
FILENAME          STRING NOT NULL,  
FILE_ROW_SEQ      NUMBER NOT NULL,  
LDTS              STRING NOT NULL,  
RSCR              STRING NOT NULL);
```

```
//No records yet entered so query produces no results  
SELECT * FROM STG_ORDER;
```

```
//STAGING AREA SETUP - STREAMS
```

```
//Create streams on the staging tables in order to easily detect and  
--incrementally process the new portion of data  
CREATE OR REPLACE STREAM STG_CUSTOMER_STRM ON TABLE  
STG_CUSTOMER;  
CREATE OR REPLACE STREAM STG_ORDER_STRM ON TABLE STG_ORDER;
```

```
//STAGING AREA SETUP - SAMPLE DATA
```

```
//We'll be producing sample data by unloading a subset of data from the TPC-H  
sample dataset then use Snowpipe to load it back  
--into the Data Vault tutorial, simulating the streaming feed.
```

```
//Every Snowflake account provides access to sample data sets. You can find  
corresponding schemas in SNOWFLAKE_SAMPLE_DATA  
--database in your object explorer. For this tutorial, we are going to use a subset of  
objects from TPC-H set, representing  
--customer and their order. We are also going to take some reference data about  
nations and regions.
```

```
//Create two stages, one for each data class type - order and customer data.  
//Note: In production environment, these would likely be internal or external stages.  
--Alternatively, these feeds could be sourced via a Kafka connector.  
CREATE OR REPLACE STAGE CUSTOMER_DATA FILE_FORMAT = (TYPE =  
JSON);  
CREATE OR REPLACE STAGE ORDER_DATA FILE_FORMAT = (TYPE = CSV) ;
```

```
//Generate and unload sample data.
```

//Use object_construct as a quick way to create an object or document from all columns
--and subsets of rows for the customer data and then offload it into CUSTOMER_DATA stage.

//ORDER data would be extracted into compressed CSV files.

//There are many additional options in COPY INTO stage construct but we are using INCLUDE_QUERY_ID
--to make it easier to generate new incremental files as we are going to run these commands
--over and over again, without a need to deal with file overloading.

//Using relatively small number of records (i.e., LIMIT) as an example in this tutorial

```
//Customer Data
COPY INTO @CUSTOMER_DATA
FROM
  (SELECT OBJECT_CONSTRUCT(*)
    FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF10.CUSTOMER LIMIT 20000)
  INCLUDE_QUERY_ID=TRUE;
```

```
//Validate
SELECT * FROM @CUSTOMER_DATA;
```

```
//Validate
SELECT METADATA$FILENAME,$1 FROM @CUSTOMER_DATA;
```

```
//Validate
LIST @CUSTOMER_DATA;
```

```
//Order Data
COPY INTO @ORDER_DATA
FROM
  (SELECT *
    FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF10.ORDERS LIMIT 50000)
  INCLUDE_QUERY_ID=TRUE;
```

```
//Validate
//Expected error from the next SQL statement
//SELECT * FROM @ORDER_DATA;
```

```
//Validate
SELECT METADATA$FILENAME,$1 FROM @ORDER_DATA;
```

```
//Validate  
LIST @ORDER_DATA;
```

//STAGING AREA SETUP - SAMPLE DATA

//Setup Snowpipe to load data from files in a stage into staging tables.

//Note: The REFRESH functionality is intended for short term use to resolve specific issues

--when Snowpipe fails to load a subset of files and is not intended for regular use.

//We're using the REFRESH in this tutorial to trigger Snowpipe explicitly to scan for new files.

--In production environment, you'll likely enable AUTO_INGEST, connecting it with your

--cloud storage events (like AWS SNS) and process new files automatically,

//Create pipe for order

```
CREATE OR REPLACE PIPE STG_ORDER_PP
```

```
AS
```

```
COPY INTO STG_ORDER
```

```
FROM
```

```
(SELECT $1,$2,$3,$4,$5,$6,$7,$8,$9,  
METADATA$FILENAME, METADATA$FILE_ROW_NUMBER,  
CURRENT_TIMESTAMP(), 'ORDER SYSTEM'  
FROM @ORDER_DATA);
```

```
ALTER PIPE STG_ORDER_PP REFRESH;
```

//Validate - no data

```
SELECT * FROM STG_ORDER;
```

```
SELECT METADATA$FILENAME,$1 FROM @ORDER_DATA;
```

//Create pipe for customer

```
CREATE OR REPLACE PIPE STG_CUSTOMER_PP
```

```
AS
```

```
COPY INTO STG_CUSTOMER
```

```
FROM
```

```
(SELECT $1,  
METADATA$FILENAME, METADATA$FILE_ROW_NUMBER,  
CURRENT_TIMESTAMP(), 'CUSTOMER SYSTEM'
```

```
FROM @CUSTOMER_DATA);
```

```
ALTER PIPE STG_CUSTOMER_PP REFRESH;
```

```
//Validate - no data
```

```
SELECT * FROM STG_CUSTOMER;
```

```
SELECT METADATA$FILENAME,$1 FROM @CUSTOMER_DATA;
```

//You should be able to see data appearing in the target tables and the stream on these tables.

//As you would expect, the number of rows in a stream is exactly the same as in the base table.

--This is because we didn't process or consume the delta of that stream yet.

//Note: it may take a few seconds before the counts are reflected in the table

```
USE SCHEMA L00_STG;
```

```
SELECT 'STG_CUSTOMER', COUNT(1) FROM STG_CUSTOMER
```

```
  UNION ALL
```

```
SELECT 'STG_ORDER', COUNT(1) FROM STG_ORDER
```

```
  UNION ALL
```

```
SELECT 'STG_ORDER_STRM', COUNT(1) FROM STG_ORDER_STRM
```

```
  UNION ALL
```

```
SELECT 'STG_CUSTOMER_STRM', COUNT(1) FROM STG_CUSTOMER_STRM;
```

//We've established the basics and now, new data will be available via stream.

//Next, we'll see if we can derive some of the business keys for the Data Vault entries in the model.

//In our example, we will model it as a view on top of the stream that should allow us to perform

--data parsing (raw_json -> columns) and business_key, hash_diff derivation on the fly.

--Another thing to notice here is the use of SHA1_binary as hashing function.

--There are many articles on choosing between MD5/SHA1(2)/other hash functions, so we won't focus on this.

//For this tutorial, we are going to use fairly common SHA1 and its BINARY version from Snowflake arsenal

--of functions that use less bytes to encode value than STRING.

```

//Create Outbound view for Customer Stream
CREATE OR REPLACE VIEW STG_CUSTOMER_STRM_OUTBOUND
AS
SELECT SRC.*,
       raw_json:C_CUSTKEY::NUMBER      C_CUSTKEY,
       raw_json:C_NAME::STRING         C_NAME,
       raw_json:C_ADDRESS::STRING      C_ADDRESS,
       raw_json:C_NATIONKEY::NUMBER     C_NATIONCODE,
       raw_json:C_PHONE::STRING         C_PHONE,
       raw_json:C_ACCTBAL::NUMBER       C_ACCTBAL,
       raw_json:C_MKTSEGMENT::STRING    C_MKTSEGMENT,
       raw_json:C_COMMENT::STRING      C_COMMENT,
-----
-- derived business key
-----
       SHA1_BINARY(UPPER(TRIM(C_CUSTKEY))) SHA1_HUB_CUSTOMER,
       SHA1_BINARY(UPPER(ARRAY_TO_STRING
           (ARRAY_CONSTRUCT
             (NVL(TRIM(C_NAME)      ,'-1'),
              NVL(TRIM(C_ADDRESS)   ,'-1'),
              NVL(TRIM(C_NATIONCODE),'-1'),
              NVL(TRIM(C_PHONE)     ,'-1'),
              NVL(TRIM(C_ACCTBAL)   ,'-1'),
              NVL(TRIM(C_MKTSEGMENT),'-1'),
              NVL(TRIM(C_COMMENT)   ,'-1')),
             '^'))))
AS CUSTOMER_HASH_DIFF
FROM STG_CUSTOMER_STRM SRC;

```

```

//Query the view to validate the results.
SELECT * FROM STG_CUSTOMER_STRM_OUTBOUND;

```

```

//Create Outbound view for Order Stream
CREATE OR REPLACE VIEW STG_ORDER_STRM_OUTBOUND
AS
SELECT SRC.*,
-----
-- derived business key
-----
       SHA1_BINARY(UPPER(TRIM(O_ORDERKEY))) SHA1_HUB_ORDER,
       SHA1_BINARY(UPPER(TRIM(O_CUSTKEY)))  SHA1_HUB_CUSTOMER,

```



```

SHA1_BINARY(UPPER(ARRAY_TO_STRING(ARRAY_CONSTRUCT
    (NVL(TRIM(O_ORDERKEY)    , '-1'),
    NVL(TRIM(O_CUSTKEY)      , '-1')),
    '^'))))
AS SHA1_LNK_CUSTOMER_ORDER,
SHA1_BINARY(UPPER(ARRAY_TO_STRING(ARRAY_CONSTRUCT
    (NVL(TRIM(O_ORDERSTATUS) , '-1'),
    NVL(TRIM(O_TOTALPRICE)   , '-1'),
    NVL(TRIM(O_ORDERDATE)    , '-1'),
    NVL(TRIM(O_ORDERPRIORITY) , '-1'),
    NVL(TRIM(O_CLERK)        , '-1'),
    NVL(TRIM(O_SHIPPRIORITY) , '-1'),
    NVL(TRIM(O_COMMENT)      , '-1')),
    '^'))))
AS ORDER_HASH_DIFF
FROM STG_ORDER_STRM SRC;

```

//Query the view to validate the results.

```
SELECT * FROM STG_ORDER_STRM_OUTBOUND;
```

//We've built out staging / inbound pipeline, ready to accommodate streaming data
and

--derived business keys that we are going to use in our Raw Data Vault

//OPTIONAL -- VALIDATE

```

USE SCHEMA L00_STG;
SELECT * FROM STG_CUSTOMER;
SELECT * FROM STG_NATION;
SELECT * FROM STG_ORDER;
SELECT * FROM STG_REGION;

```

```

SELECT * FROM STG_CUSTOMER_STRM_OUTBOUND;
SELECT * FROM STG_ORDER_STRM_OUTBOUND;

```

```

SELECT * FROM @CUSTOMER_DATA;
//SELECT * FROM @ORDER_DATA;

```

```

SELECT * FROM STG_CUSTOMER_STRM;
SELECT * FROM STG_ORDER_STRM;

```


//BUILD RAW DATA VAULT - CUSTOMER AND ORDER HUBS

//Deploy DDL for the HUBs, LINKs, and SATELLITE tables.
USE SCHEMA L10_RDV;

//Create Customer Hub

```
CREATE OR REPLACE TABLE HUB_CUSTOMER
(SHA1_HUB_CUSTOMER  BINARY  NOT NULL,
 C_CUSTKEY          NUMBER  NOT NULL,
 LDTS              TIMESTAMP NOT NULL,
 RSCR              STRING   NOT NULL,
 CONSTRAINT PK_HUB_CUSTOMER PRIMARY
KEY(SHA1_HUB_CUSTOMER));
```

//Create Order Hub

```
CREATE OR REPLACE TABLE HUB_ORDER
(SHA1_HUB_ORDER     BINARY  NOT NULL,
 O_ORDERKEY         NUMBER  NOT NULL,
 LDTS              TIMESTAMP NOT NULL,
 RSCR              STRING   NOT NULL,
 CONSTRAINT PK_HUB_ORDER PRIMARY KEY(SHA1_HUB_ORDER));
```

//BUILD RAW DATA VAULT - CUSTOMER AND ORDER SATELLITES

//Create Customer Satellite

```
CREATE OR REPLACE TABLE SAT_CUSTOMER
(SHA1_HUB_CUSTOMER  BINARY  NOT NULL,
 LDTS              TIMESTAMP NOT NULL,
 C_NAME            STRING,
 C_ADDRESS         STRING,
 C_PHONE          STRING,
 C_ACCTBAL         NUMBER,
 C_MKTSEGMENT      STRING,
 C_COMMENT         STRING,
 NATIONCODE        NUMBER,
 HASH_DIFF         BINARY  NOT NULL,
 RSCR              STRING   NOT NULL,
```

```
    CONSTRAINT PK_SAT_CUSTOMER PRIMARY
KEY(SHA1_HUB_CUSTOMER, LDTS),
    CONSTRAINT FK_SAT_CUSTOMER FOREIGN
KEY(SHA1_HUB_CUSTOMER) REFERENCES HUB_CUSTOMER);
```

//Create ORDER Satellite

```
CREATE OR REPLACE TABLE SAT_ORDER
(SHA1_HUB_ORDER BINARY NOT NULL,
LDTS TIMESTAMP NOT NULL,
O_ORDERSTATUS STRING,
O_TOTALPRICE NUMBER,
O_ORDERDATE DATE,
O_ORDERPRIORITY STRING,
O_CLERK STRING,
O_SHIPPRIORITY NUMBER,
O_COMMENT STRING,
HASH_DIFF BINARY NOT NULL,
RSCR STRING NOT NULL,
CONSTRAINT PK_SAT_ORDER PRIMARY KEY(SHA1_HUB_ORDER, LDTS),
CONSTRAINT FK_SAT_ORDER FOREIGN KEY(SHA1_HUB_ORDER)
REFERENCES HUB_ORDER);
```

//BUILD RAW DATA VAULT - CUSTOMER_ORDER Link

//Create the CUSTOMER_ORDER Link

```
CREATE OR REPLACE TABLE LNK_CUSTOMER_ORDER
(SHA1_LNK_CUSTOMER_ORDER BINARY NOT NULL,
SHA1_HUB_CUSTOMER BINARY,
SHA1_HUB_ORDER BINARY,
LDTS TIMESTAMP NOT NULL,
RSCR STRING NOT NULL,
CONSTRAINT PK_LNK_CUSTOMER_ORDER PRIMARY
KEY(SHA1_LNK_CUSTOMER_ORDER),
CONSTRAINT FK1_LNK_CUSTOMER_ORDER FOREIGN
KEY(SHA1_HUB_CUSTOMER) REFERENCES HUB_CUSTOMER,
CONSTRAINT FK2_LNK_CUSTOMER_ORDER FOREIGN
KEY(SHA1_HUB_ORDER) REFERENCES HUB_ORDER);
```

//BUILD RAW DATA VAULT - REFERENCE TABLES

//Create Region Reference

```
CREATE OR REPLACE TABLE REF_REGION
```

```

(REGIONCODE      NUMBER,
 LDTS            TIMESTAMP,
 RSCR            STRING NOT NULL,
 R_NAME          STRING,
 R_COMMENT        STRING,
 CONSTRAINT PK_REF_REGION PRIMARY KEY (REGIONCODE))
AS
SELECT R_REGIONKEY, LDTS, RSCR, R_NAME, R_COMMENT
FROM L00_STG.STG_REGION;

```

```

//Validate
SELECT * FROM REF_REGION;

```

```

//Create Nation Reference
CREATE OR REPLACE TABLE REF_NATION
(
  NATIONCODE      NUMBER,
  REGIONCODE      NUMBER,
  LDTS            TIMESTAMP,
  RSCR            STRING NOT NULL,
  N_NAME          STRING,
  N_COMMENT        STRING,
  CONSTRAINT PK_REF_NATION PRIMARY KEY (NATIONCODE),
  CONSTRAINT FK_REF_REGION FOREIGN KEY (REGIONCODE)
REFERENCES REF_REGION(REGIONCODE))
AS
SELECT N_NATIONKEY, N_REGIONKEY, LDTS, RSCR, N_NAME,
N_COMMENT
FROM L00_STG.STG_NATION;

```

```

//Validate
SELECT * FROM REF_NATION;

```

//BUILD RAW DATA VAULT - CREATE TASKS

//Now we have source data waiting in our staging streams and views and we have target Raw Data Vaults tables.

--We now need to connect the dots. We are going to create tasks, one per each stream so whenever

--there is new records coming in a stream, that delta will be incrementally propagated to all

--dependent RDV models in one go.

//To achieve that, we are going to use multi-table insert functionality

--As you can see, tasks can be set up to run on a

--pre-defined frequency (every 1 minute in our example) and use dedicated virtual warehouse

--as a compute power (in our tutorial we are going to use same warehouse for all tasks,

--though this could be as granular as needed).

--Also, before waking up compute resource, tasks are going to check that there is data

--in a corresponding stream to process.

//Create Task for Customer Stream

CREATE OR REPLACE TASK CUSTOMER_STRM_TSK

WAREHOUSE = DV_RDV_WH

SCHEDULE = '1 minute'

WHEN

SYSTEM\$STREAM_HAS_DATA('L00_STG.STG_CUSTOMER_STRM')

AS INSERT ALL

WHEN (SELECT COUNT(1) FROM HUB_CUSTOMER TGT WHERE
TGT.SHA1_HUB_CUSTOMER = SRC_SHA1_HUB_CUSTOMER) = 0

THEN INTO HUB_CUSTOMER

(SHA1_HUB_CUSTOMER, C_CUSTKEY, LDTS, RSCR)

VALUES (SRC_SHA1_HUB_CUSTOMER, SRC_C_CUSTKEY, SRC_LDTS,
SRC_RSCR)

WHEN (SELECT COUNT(1) FROM SAT_CUSTOMER TGT WHERE
TGT.SHA1_HUB_CUSTOMER = SRC_SHA1_HUB_CUSTOMER AND
TGT.HASH_DIFF = SRC_CUSTOMER_HASH_DIFF) = 0

THEN INTO SAT_CUSTOMER

(SHA1_HUB_CUSTOMER, LDTS, C_NAME, C_ADDRESS, C_PHONE,
C_ACCTBAL, C_MKTSEGMENT, C_COMMENT,
NATIONCODE, HASH_DIFF, RSCR)

VALUES (SRC_SHA1_HUB_CUSTOMER, SRC_LDTS, SRC_C_NAME,
SRC_C_ADDRESS, SRC_C_PHONE, SRC_C_ACCTBAL,
SRC_C_MKTSEGMENT, SRC_C_COMMENT, SRC_NATIONCODE,
SRC_CUSTOMER_HASH_DIFF, SRC_RSCR)

SELECT

SHA1_HUB_CUSTOMER SRC_SHA1_HUB_CUSTOMER,

C_CUSTKEY SRC_C_CUSTKEY,

C_NAME SRC_C_NAME,

C_ADDRESS SRC_C_ADDRESS,

C_NATIONCODE SRC_NATIONCODE,

C_PHONE SRC_C_PHONE,

```

C_ACCTBAL      SRC_C_ACCTBAL,
C_MKTSEGMENT   SRC_C_MKTSEGMENT,
C_COMMENT      SRC_C_COMMENT,
CUSTOMER_HASH_DIFF SRC_CUSTOMER_HASH_DIFF,
LDTS           SRC_LDTS,
RSCR           SRC_RSCR
FROM L00_STG.STG_CUSTOMER_STRM_OUTBOUND SRC;

```

```

ALTER TASK CUSTOMER_STRM_TSK RESUME;

```

```

//Create task for ORDER Stream

```

```

CREATE OR REPLACE TASK ORDER_STRM_TSK
  WAREHOUSE = DV_RDV_WH
  SCHEDULE = '1 minute'
WHEN
  SYSTEM$STREAM_HAS_DATA('L00_STG.STG_ORDER_STRM')
AS INSERT ALL
WHEN (SELECT COUNT(1) FROM HUB_ORDER TGT WHERE
TGT.SHA1_HUB_ORDER = SRC_SHA1_HUB_ORDER) = 0
THEN INTO HUB_ORDER
  (SHA1_HUB_ORDER, O_ORDERKEY, LDTS, RSCR)
  VALUES (SRC_SHA1_HUB_ORDER, SRC_O_ORDERKEY, SRC_LDTS,
SRC_RSCR)
WHEN (SELECT COUNT(1) FROM SAT_ORDER TGT WHERE
TGT.SHA1_HUB_ORDER = SRC_SHA1_HUB_ORDER AND TGT.HASH_DIFF =
SRC_ORDER_HASH_DIFF) = 0
THEN INTO SAT_ORDER
  (SHA1_HUB_ORDER, LDTS, O_ORDERSTATUS, O_TOTALPRICE,
O_ORDERDATE, O_ORDERPRIORITY, O_CLERK,
  O_SHIPPRIORITY, O_COMMENT, HASH_DIFF, RSCR)
  VALUES (SRC_SHA1_HUB_ORDER, SRC_LDTS, SRC_O_ORDERSTATUS,
SRC_O_TOTALPRICE, SRC_O_ORDERDATE,
  SRC_O_ORDERPRIORITY, SRC_O_CLERK, SRC_O_SHIPPRIORITY,
SRC_O_COMMENT, SRC_ORDER_HASH_DIFF, SRC_RSCR)
WHEN (SELECT COUNT(1) FROM LNK_CUSTOMER_ORDER TGT
  WHERE TGT.SHA1_LNK_CUSTOMER_ORDER =
SRC_SHA1_LNK_CUSTOMER_ORDER) = 0
THEN INTO LNK_CUSTOMER_ORDER
  (SHA1_LNK_CUSTOMER_ORDER, SHA1_HUB_CUSTOMER,
SHA1_HUB_ORDER, LDTS, RSCR)

```

```

VALUES (SRC_SHA1_LNK_CUSTOMER_ORDER,
SRC_SHA1_HUB_CUSTOMER, SRC_SHA1_HUB_ORDER, SRC_LDTS,
SRC_RSCR)
SELECT
  SHA1_HUB_ORDER      SRC_SHA1_HUB_ORDER,
  SHA1_LNK_CUSTOMER_ORDER SRC_SHA1_LNK_CUSTOMER_ORDER,
  SHA1_HUB_CUSTOMER    SRC_SHA1_HUB_CUSTOMER,
  O_ORDERKEY           SRC_O_ORDERKEY,
  O_ORDERSTATUS        SRC_O_ORDERSTATUS,
  O_TOTALPRICE         SRC_O_TOTALPRICE,
  O_ORDERDATE          SRC_O_ORDERDATE,
  O_ORDERPRIORITY      SRC_O_ORDERPRIORITY,
  O_CLERK              SRC_O_CLERK,
  O_SHIPPRIORITY        SRC_O_SHIPPRIORITY,
  O_COMMENT            SRC_O_COMMENT,
  ORDER_HASH_DIFF      SRC_ORDER_HASH_DIFF,
  LDTS                 SRC_LDTS,
  RSCR                 SRC_RSCR
FROM L00_STG.STG_ORDER_STRM_OUTBOUND SRC;

ALTER TASK ORDER_STRM_TSK  RESUME;

```

//Once tasks are created and resumed (by default, they are initially suspended),
//let's have a look on the task execution history to see how the process will start.

```

SELECT *
FROM TABLE(INFORMATION_SCHEMA.TASK_HISTORY())
ORDER BY SCHEDULED_TIME DESC;

```

//Notice how after successful execution, next two tasks run were automatically
SKIPPED as there were nothing in the stream and there is nothing to do.

//We can also check content and stats of the objects involved.

--Please notice that views on streams in our staging area are no longer returning any
rows.

--This is because that delta of changes was consumed by a successfully completed
DML transaction (in our case, embedded in tasks).

--This way you don't need to spend any time implementing incremental detection /
processing logic on the application side.

```

SELECT 'HUB_CUSTOMER', COUNT(1) FROM HUB_CUSTOMER
UNION ALL
SELECT 'HUB_ORDER', COUNT(1) FROM HUB_ORDER

```

```

UNION ALL
SELECT 'SAT_CUSTOMER', COUNT(1) FROM SAT_CUSTOMER
UNION ALL
SELECT 'SAT_ORDER', COUNT(1) FROM SAT_ORDER
UNION ALL
SELECT 'LNK_CUSTOMER_ORDER', COUNT(1) FROM
LNK_CUSTOMER_ORDER
UNION ALL
SELECT 'L00_STG.STG_CUSTOMER_STRM_OUTBOUND', COUNT(1) FROM
L00_STG.STG_CUSTOMER_STRM_OUTBOUND
UNION ALL
SELECT 'L00_STG.STG_ORDER_STRM_OUTBOUND', COUNT(1) FROM
L00_STG.STG_ORDER_STRM_OUTBOUND;

```

//We now have data in our Raw Data Vault core structures. Let's move on and talk about the concept of virtualization for building your near-real time Data Vault solution.

```

-----
-----
-----

```

//VIEWS FOR AGILE REPORTING

//One of the great benefits of having the compute power from Snowflake is that now it is totally possible to have most of your business vault and information marts in a Data Vault architecture be built exclusively from views. There is no longer a need to have the argument that there are "too many joins" or that the response won't be fast enough. The elasticity of the Snowflake virtual warehouses combined with our dynamic optimization engine have solved that problem. (For more details, see this post)

//If you really want to deliver data to the business users and data scientists in NRT, in our opinion using views is the only option. Once you have the streaming loads built to feed your Data Vault, the fastest way to make that data visible downstream will be views. Using views allows you to deliver the data faster by eliminating any latency that would be incurred by having additional ELT processes between the Data Vault and the data consumers downstream.

//All the business logic, alignment, and formatting of the data can be in the view code. That means fewer moving parts to debug, and reduces the storage needed as well.

//Looking at the diagram above you will see an example of how virtualization could fit in the architecture. Here, solid lines are representing physical tables and dotted lines - views. You incrementally ingest data into Raw Data Vault and all downstream

transformations are applied as views. From a data consumer perspective when working with a virtualized information mart, the query always shows everything known by your data vault, right up to the point the query was submitted.

//With Snowflake you have the ability to provide as much compute as required, on-demand, without a risk of causing performance impact on any surrounding processes and pay only for what you use. This makes materialization of transformations in layers like Business Data Vault and Information delivery an option rather than a must-have. Instead of "optimizing upfront" you can now make this decision based on the usage pattern characteristics, such as frequency of use, type of queries, latency requirements, readiness of the requirements etc.

//Many modern data engineering automation frameworks are already actively supporting virtualization of logic. Several tools offer a low-code or configuration-like ability to switch between materializing an object as a view or a physical table, automatically generating all required DDL & DML. This could be applied on specific objects, layers or/and be environment specific. So even if you start with a view, you can easily refactor to use a table if user requirements evolve.

//As said before, virtualization is not only a way to improve time-to-value and provide near real time access to the data, given the scalability and workload isolation of Snowflake, virtualization also is a design technique that could make your Data Vault excel: minimizing cost-of-change, accelerating the time-to-delivery and becoming an extremely agile, future proof solution for ever growing business needs.

//BUILD: Business Data Vault

//As a quick example of using views for transformations we just discussed, here is how enrichment of

--customer descriptive data could happen in Business Data Vault, connecting data received from

--source with some reference data.

//Let's create a view that will perform these additional derivations on the fly.

//Assuming non-functional capabilities are satisfying our requirements,

--deploying (and re-deploying a new version) transformations in this way is super easy.

USE SCHEMA L20_BDV;

CREATE OR REPLACE VIEW SAT_CUSTOMER_BV

AS

```

SELECT RSC.SHA1_HUB_CUSTOMER, RSC.LDTS, RSC.C_NAME,
RSC.C_ADDRESS, RSC.C_PHONE,
    RSC.C_ACCTBAL, RSC.C_MKTSEGMENT, RSC.C_COMMENT,
RSC.NATIONCODE, RSC.RSCR,
    -- derived
    RRN.N_NAME    NATION_NAME,
    RRR.R_NAME    REGION_NAME
FROM L10_RDV.SAT_CUSTOMER RSC
    LEFT OUTER JOIN L10_RDV.REF_NATION RRN
        ON (RSC.NATIONCODE = RRN.NATIONCODE)
    LEFT OUTER JOIN L10_RDV.REF_REGION RRR
        ON (RRN.REGIONCODE = RRR.REGIONCODE);

```

//Verify

```

SELECT * FROM SAT_CUSTOMER_BV;

```

//Now let's imagine we have a heavier transformation to perform that it would make more sense

--to materialize it as a table. It could be more data volume, could be more complex logic,

--PITs, bridges or even an object that will be used frequently and by many users.

--For this case, let's first build a new business satellite that for illustration purposes

--will be deriving additional classification / tiering for orders based on the conditional logic.

```

CREATE OR REPLACE TABLE SAT_ORDER_BV
    (SHA1_HUB_ORDER    BINARY    NOT NULL,
    LDTS              TIMESTAMP NOT NULL,
    O_ORDERSTATUS     STRING,
    O_TOTALPRICE      NUMBER,
    O_ORDERDATE       DATE,
    O_ORDERPRIORITY   STRING,
    O_CLERK           STRING,
    O_SHIPPRIORITY    NUMBER,
    O_COMMENT         STRING,
    HASH_DIFF         BINARY NOT NULL,
    RSCR              STRING NOT NULL,
    -- additional attributes
    ORDER_PRIORITY_BUCKET STRING,
    CONSTRAINT PK_SAT_ORDER PRIMARY KEY(SHA1_HUB_ORDER, LDTS),
    CONSTRAINT FK_SAT_ORDER FOREIGN KEY(SHA1_HUB_ORDER)
REFERENCES L10_RDV.HUB_ORDER)
AS
SELECT SHA1_HUB_ORDER, LDTS, O_ORDERSTATUS, O_TOTALPRICE,
O_ORDERDATE, O_ORDERPRIORITY,

```

```

O_CLERK, O_SHIPPRIORITY, O_COMMENT, HASH_DIFF, RSCR,
-- derived additional attributes
CASE WHEN O_ORDERPRIORITY IN ('2-HIGH', '1-URGENT') AND
O_TOTALPRICE >= 200000 THEN 'Tier-1'
      WHEN O_ORDERPRIORITY IN ('3-MEDIUM', '2-HIGH', '1-URGENT') AND
O_TOTALPRICE BETWEEN 150000 AND 200000 THEN 'Tier-2'
      ELSE 'Tier-3'
END ORDER_PRIORITY_BUCKET
FROM L10_RDV.SAT_ORDER;

```

```

//VALIDATE
SELECT * FROM SAT_ORDER_BV;

```

//What we are going to do from processing/orchestration perspective is extending our --ORDER processing pipeline so that when the task populates a l10_rdv.sat_ORDER this

--will generate a new stream of changes and these changes are going to be propagated

--by a dependent task to l20_bdv.sat_ORDER_bv. This is super easy to do as tasks in

--Snowflake can be not only schedule-based but also start automatically once the --parent task is completed.

```

CREATE OR REPLACE STREAM L10_RDV.SAT_ORDER_STRM ON TABLE
L10_RDV.SAT_ORDER;

```

```

ALTER TASK L10_RDV.ORDER_STRM_TSK SUSPEND;

```

```

CREATE OR REPLACE TASK
L10_RDV.HUB_ORDER_STRM_SAT_ORDER_BV_TSK
WAREHOUSE = DV_RDV_WH
AFTER L10_RDV.ORDER_STRM_TSK
AS
INSERT INTO L20_BDV.SAT_ORDER_BV
SELECT
  SHA1_HUB_ORDER, LDTS, O_ORDERSTATUS, O_TOTALPRICE,
O_ORDERDATE, O_ORDERPRIORITY,
  O_CLERK, O_SHIPPRIORITY, O_COMMENT, HASH_DIFF, RSCR,
  -- derived additional attributes
CASE
  WHEN O_ORDERPRIORITY IN ('2-HIGH', '1-URGENT') AND O_TOTALPRICE
>= 200000 THEN 'Tier-1'

```

```
        WHEN O_ORDERPRIORITY IN ('3-MEDIUM', '2-HIGH', '1-URGENT') AND  
O_TOTALPRICE BETWEEN 150000 AND 200000 THEN 'Tier-2'  
        ELSE 'Tier-3'  
    END ORDER_PRIORITY_BUCKET  
FROM SAT_ORDER_STRM;
```

```
ALTER TASK L10_RDV.HUB_ORDER_STRM_SAT_ORDER_BV_TSK RESUME;  
ALTER TASK L10_RDV.ORDER_STRM_TSK RESUME;
```

//Now let's go back to our staging area to process another slice of data to test the task

```
USE SCHEMA L00_STG;  
COPY INTO @ORDER_DATA  
FROM  
    (SELECT * FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF10.ORDERS LIMIT  
125000)  
--Removed the limit  
INCLUDE_QUERY_ID=TRUE;
```

```
ALTER PIPE STG_ORDER_PP REFRESH;
```

//Data is not automatically flowing through all the layers via asynchronous tasks.

--With the results, you can validate:

```
SELECT 'L00_STG.STG_ORDER', COUNT(1) FROM L00_STG.STG_ORDER  
UNION ALL  
SELECT 'L00_STG.STG_ORDER_STRM', COUNT(1) FROM  
L00_STG.STG_ORDER_STRM  
UNION ALL  
SELECT 'L10_RDV.SAT_ORDER', COUNT(1) FROM L10_RDV.SAT_ORDER  
UNION ALL  
SELECT 'L10_RDV.SAT_ORDER_STRM', COUNT(1) FROM  
L10_RDV.SAT_ORDER_STRM  
UNION ALL  
SELECT 'L20_BDV.SAT_ORDER_BV', COUNT(1) FROM  
L20_BDV.SAT_ORDER_BV;
```

```
SELECT *  
FROM TABLE(INFORMATION_SCHEMA.TASK_HISTORY())  
ORDER BY SCHEDULED_TIME DESC;
```

//Now let's go back to our staging area to process another slice of data to test the task

```
USE SCHEMA L00_STG;  
COPY INTO @ORDER_DATA  
FROM  
    (SELECT * FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF10.ORDERS)  
--Removed the limit  
INCLUDE_QUERY_ID=TRUE;
```

```
ALTER PIPE STG_ORDER_PP REFRESH;
```

//Data is not automatically flowing through all the layers via asynchronous tasks.

--With the results, you can validate:

```
SELECT 'L00_STG.STG_ORDER', COUNT(1) FROM L00_STG.STG_ORDER  
UNION ALL  
SELECT 'L00_STG.STG_ORDER_STRM', COUNT(1) FROM  
L00_STG.STG_ORDER_STRM  
UNION ALL  
SELECT 'L10_RDV.SAT_ORDER', COUNT(1) FROM L10_RDV.SAT_ORDER  
UNION ALL  
SELECT 'L10_RDV.SAT_ORDER_STRM', COUNT(1) FROM  
L10_RDV.SAT_ORDER_STRM  
UNION ALL  
SELECT 'L20_BDV.SAT_ORDER_BV', COUNT(1) FROM  
L20_BDV.SAT_ORDER_BV;
```

```
SELECT *  
FROM TABLE(INFORMATION_SCHEMA.TASK_HISTORY())  
ORDER BY SCHEDULED_TIME DESC;
```

//Now let's go back to our staging area to process another slice of data to test the task

```
USE SCHEMA L00_STG;  
COPY INTO @ORDER_DATA  
FROM  
    (SELECT * FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF10.ORDERS LIMIT  
33333)
```

```
INCLUDE_QUERY_ID=TRUE;
```

```
ALTER PIPE STG_ORDER_PP REFRESH;
```

```
//Data is not automatically flowing through all the layers via asynchronous tasks.
```

```
--With the results, you can validate:
```

```
SELECT 'L00_STG.STG_ORDER', COUNT(1) FROM L00_STG.STG_ORDER  
UNION ALL
```

```
SELECT 'L00_STG.STG_ORDER_STRM', COUNT(1) FROM  
L00_STG.STG_ORDER_STRM  
UNION ALL
```

```
SELECT 'L10_RDV.SAT_ORDER', COUNT(1) FROM L10_RDV.SAT_ORDER  
UNION ALL
```

```
SELECT 'L10_RDV.SAT_ORDER_STRM', COUNT(1) FROM  
L10_RDV.SAT_ORDER_STRM  
UNION ALL
```

```
SELECT 'L20_BDV.SAT_ORDER_BV', COUNT(1) FROM  
L20_BDV.SAT_ORDER_BV;
```

```
SELECT *  
FROM TABLE(INFORMATION_SCHEMA.TASK_HISTORY())  
ORDER BY SCHEDULED_TIME DESC;
```

```
//Now let's go back to our staging area to process another slice of data to test the  
task
```

```
USE SCHEMA L00_STG;
```

```
COPY INTO @ORDER_DATA
```

```
FROM
```

```
(SELECT * FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF10.ORDERS)
```

```
--Removed the limit
```

```
INCLUDE_QUERY_ID=TRUE;
```

```
ALTER PIPE STG_ORDER_PP REFRESH;
```

```
//Data is not automatically flowing through all the layers via asynchronous tasks.
```

--With the results, you can validate:

```
SELECT 'L00_STG.STG_ORDER', COUNT(1) FROM L00_STG.STG_ORDER
UNION ALL
SELECT 'L00_STG.STG_ORDER_STRM', COUNT(1) FROM
L00_STG.STG_ORDER_STRM
UNION ALL
SELECT 'L10_RDV.SAT_ORDER', COUNT(1) FROM L10_RDV.SAT_ORDER
UNION ALL
SELECT 'L10_RDV.SAT_ORDER_STRM', COUNT(1) FROM
L10_RDV.SAT_ORDER_STRM
UNION ALL
SELECT 'L20_BDV.SAT_ORDER_BV', COUNT(1) FROM
L20_BDV.SAT_ORDER_BV;
```

```
SELECT *
FROM TABLE(INFORMATION_SCHEMA.TASK_HISTORY())
ORDER BY SCHEDULED_TIME DESC;
```


//BUILD: Information Delivery

//When it comes to the Information Delivery layer we are not changing the meaning of data,

--but we may change format to simplify users to access and work with the

--data products/output interfaces. Different consumers may have different needs and preferences,

--some would prefer star/snowflake dimensional schemas, some would adhere to use

--flattened objects or even transform data into JSON/parquet objects.

//First things we would like to add to simplify working with satellites

--is creating views that shows latest version for each key.

-- RDV curr views

```
USE SCHEMA L10_RDV;
```

```
CREATE OR REPLACE VIEW SAT_CUSTOMER_CURR_VW
AS
  SELECT * FROM SAT_CUSTOMER
QUALIFY LEAD(LDTS) OVER (PARTITION BY SHA1_HUB_CUSTOMER ORDER
BY LDTS) IS NULL;
```

```
CREATE OR REPLACE VIEW SAT_ORDER_CURR_VW
AS
  SELECT * FROM SAT_ORDER
QUALIFY LEAD(LDTS) OVER (PARTITION BY SHA1_HUB_ORDER ORDER BY
LDTS) IS NULL;
```

```
-----
-- BDV curr views
-----
```

```
USE SCHEMA L20_BDV;
```

```
CREATE OR REPLACE VIEW SAT_ORDER_BV_CURR_VW
AS
  SELECT * FROM SAT_ORDER_BV
QUALIFY LEAD(LDTS) OVER (PARTITION BY SHA1_HUB_ORDER ORDER BY
LDTS) IS NULL;
```

```
CREATE VIEW SAT_CUSTOMER_BV_CURR_VW
AS
  SELECT * FROM SAT_CUSTOMER_BV
QUALIFY LEAD(LDTS) OVER (PARTITION BY SHA1_HUB_CUSTOMER ORDER
BY LDTS) IS NULL;
```

//Let's create a simple dimensional structure. Again, we will keep it virtual(as views) to start with,

--but you already know that depending on access characteristics required any of these

--could be selectively materialized.

```
USE SCHEMA L30_INFO;
```

```
-- DIM TYPE 1
```

```
CREATE OR REPLACE VIEW DIM1_CUSTOMER
```

```
AS
```

```
SELECT
```

```
  HUB.SHA1_HUB_CUSTOMER
```

```
          AS DIM_CUSTOMER_KEY,
```



```

    SAT.LDTS                AS EFFECTIVE_DTS,
    HUB.C_CUSTKEY            AS CUSTOMER_ID,
    SAT.RSCR                AS RECORD_SOURCE,
    SAT.*
FROM
    L10_RDV.HUB_CUSTOMER    HUB,
    L20_BDV.SAT_CUSTOMER_BV_CURR_VW    SAT
WHERE HUB.SHA1_HUB_CUSTOMER = SAT.SHA1_HUB_CUSTOMER;

```

```

-- DIM TYPE 1
CREATE OR REPLACE VIEW DIM1_ORDER
AS
SELECT

```

```

    HUB.SHA1_HUB_ORDER      AS DIM_ORDER_KEY,
    SAT.LDTS                AS EFFECTIVE_DTS,
    HUB.O_ORDERKEY          AS ORDER_ID,
    SAT.RSCR                AS RECORD_SOURCE,
    SAT.*
FROM
    L10_RDV.HUB_ORDER      HUB,
    L20_BDV.SAT_ORDER_BV_CURR_VW    SAT
WHERE HUB.SHA1_HUB_ORDER = SAT.SHA1_HUB_ORDER;

```

```

-- FACT table
CREATE OR REPLACE VIEW FCT_CUSTOMER_ORDER
AS
SELECT
    LNK.LDTS                AS EFFECTIVE_DTS,
    LNK.RSCR                AS RECORD_SOURCE,
    LNK.SHA1_HUB_CUSTOMER   AS DIM_CUSTOMER_KEY,
    LNK.SHA1_HUB_ORDER      AS DIM_ORDER_KEY
-- this is a factless fact, but here you can add any measures, calculated or derived
FROM L10_RDV.LNK_CUSTOMER_ORDER    LNK;

```

//All good so far?

//Now lets try to query fct_customer_order. You may find that the view does not return any rows. Why?

--If you remember, when we were unloading sample data, we took a subset of random orders and a subset of random customers.

--Thus, it is possible that there won't be any overlap, Therefore doing the inner join with dim1_order will likely result

--in all rows being eliminated from the resultset. Thankfully we are using Data Vault and all we need to do is go and load
--the full customer dataset. Just think about it, there is no need to reprocess any links or fact tables simply because
--customer/reference feed was incomplete. Lets go and see if we can resolve this.

```
USE SCHEMA L00_STG;  
COPY INTO @CUSTOMER_DATA  
FROM  
    (SELECT OBJECT_CONSTRUCT(*)  
      FROM SNOWFLAKE_SAMPLE_DATA.TPCH_SF10.CUSTOMER)  
    -- removed LIMIT  
INCLUDE_QUERY_ID=TRUE;
```

```
ALTER PIPE STG_CUSTOMER_PP REFRESH;
```

//All you need to do now is just wait a few seconds whilst our continuous data pipeline will automatically propagate
--new customer data into Raw Data Vault. Quick check for the records count in customer dimension now shows
--that there are 1.5Mn records:

```
USE SCHEMA L30_INFO;  
SELECT COUNT(1) FROM DIM1_CUSTOMER;  
//Note that if the result is zero, wait a few seconds and rerun the query
```

//Finally lets wear user's hat and run a query to break down ORDER by nation, region and ORDER_PRIORITY_BUCKET which
-- are all attributes we derived in Business Data Vault. As we are using Snowsight, why not quickly creating a chart
--from this result set to better understand the data. For this simply click on the 'Chart' section on the bottom pane
--and put attributes/measures as it is shown on the screenshot below.

```
SELECT DC.NATION_NAME, DC.REGION_NAME,  
DO.ORDER_PRIORITY_BUCKET, COUNT(1) CNT_ORDER  
FROM FCT_CUSTOMER_ORDER FCT,  
     DIM1_CUSTOMER      DC,  
     DIM1_ORDER          DO  
WHERE FCT.DIM_CUSTOMER_KEY = DC.DIM_CUSTOMER_KEY AND  
      FCT.DIM_ORDER_KEY = DO.DIM_ORDER_KEY  
GROUP BY 1,2,3;
```

```
SELECT * FROM FCT_CUSTOMER_ORDER;  
SELECT * FROM DIM1_CUSTOMER;  
SELECT * FROM DIM1_ORDER;
```

//Simplicity of engineering, openness, scalable performance, enterprise-grade governance enabled by the core of the Snowflake platform are now allowing teams to focus on what matters most for the business and build truly agile, collaborative data environments. Teams can now connect data from all parts of the landscape, until there are no stones left unturned. They are even tapping into new datasets via live access to the Snowflake Data Marketplace. The Snowflake Data Cloud combined with a Data Vault 2.0 approach is allowing teams to democratize access to all their data assets at any scale. We can now easily derive more and more value through insights and intelligence, day after day, bringing businesses to the next level of being truly data-driven.

//Delivering more usable data faster is no longer an option for today's business environment. Using the Snowflake platform, combined with the Data Vault 2.0 architecture it is now possible to build a world class analytics platform that delivers data for all users in near real-time.

//Option to suspend tasks

```
ALTER TASK L10_RDV.ORDER_STRM_TSK SUSPEND;  
ALTER TASK L10_RDV.CUSTOMER_STRM_TSK SUSPEND;  
ALTER TASK L10_RDV.HUB_ORDER_STRM_SAT_ORDER_BV_TSK SUSPEND;
```

//Cleanup

```
USE WAREHOUSE COMPUTE_WH;  
DROP DATABASE DV_TUTORIAL;  
DROP WAREHOUSE DV_GENERIC_WH;  
DROP WAREHOUSE DV_RDV_WH;
```